

# RU File System using FUSE

## Project 4

**Authors:** Krishna Sathyamurthy (ks2025), Abhinav Bharadwaj Sarathy (ab2812)

## Project 4

### Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Implementation of File System</b>	<b>2</b>
2.1	Disk . . . . .	2
2.1.1	dev_init . . . . .	2
2.1.2	dev_open . . . . .	2
2.1.3	dev_close . . . . .	2
2.1.4	bio_read . . . . .	2
2.1.5	bio_write . . . . .	2
2.2	Inode . . . . .	2
2.2.1	readi . . . . .	3
2.2.2	writei . . . . .	3
2.3	Bitmap . . . . .	3
2.3.1	get_bitmap . . . . .	3
2.3.2	set_bitmap . . . . .	3
2.3.3	unset_bitmap . . . . .	3
2.3.4	get_avail_ino . . . . .	3
2.3.5	get_avail_blkno . . . . .	3
2.4	Files and Directories . . . . .	3
2.4.1	dir_find . . . . .	3
2.4.2	dir_add . . . . .	3
2.4.3	get_node_by_path . . . . .	4
2.4.4	dir_remove . . . . .	4
2.5	FUSE handlers . . . . .	4
2.5.1	rufs_init . . . . .	4
2.5.2	rufs_mkfs . . . . .	4
2.5.3	rufs_destroy . . . . .	4
2.5.4	rufs_getattr . . . . .	4
2.5.5	rufs_opendir . . . . .	4
2.5.6	rufs_readdir . . . . .	4
2.5.7	rufs_mkdir . . . . .	4
2.5.8	rufs_create . . . . .	4
2.5.9	rufs_open . . . . .	4
2.5.10	rufs_read . . . . .	5
2.5.11	rufs_write . . . . .	5
2.5.12	rufs_rmdir . . . . .	5
2.5.13	rufs_unlink . . . . .	5
<b>3</b>	<b>Benchmark testing</b>	<b>6</b>
<b>4</b>	<b>Challenges faced while implementing</b>	<b>7</b>
4.1	Bitmaps . . . . .	7
4.2	Indirect pointers . . . . .	7

## 1 Introduction

For this project, we'll be implementing a RU File System (RUFFS) on top of the FUSE library. We are given with the following -

```
// A supreme structure which handles all the underlying data
struct superblock;

// Used to store all the directory entries in the file system
struct dirent;

// Used to store the meta data of files/directories in the system
struct inode;
```

Before jumping to the implementation, let's brief what we'll be doing here. First thing which is essential for a file system is the disk. We'll be seeing how we have implemented the disk as blocks, how the blocks are managed efficiently. Further, we'll see how files and directories are implemented to manage. And lastly how FUSE handlers are operated to support system calls.

## 2 Implementation of File System

### 2.1 Disk

The Disk implementation is the foundation to our file system. We have considered the disk to be a flat file of DISK\_SIZE (32MB) bytes. Further we have partitioned the disks as blocks of BLOCK\_SIZE (4KB). By doing this, it becomes fairly simple to manage the entities in the file system.

#### 2.1.1 dev\_init

This function takes the diskfile\_path as input; which is assigned in the main function of 'ruffs.c'. It creates a flat file of size DISK\_SIZE bytes (32MB), which acts as our disk for the file system. In case of failure in opening the diskfile, we exit the program with error code EXIT\_FAILURE.

#### 2.1.2 dev\_open

This function takes the char pointer diskfile\_path as input and opens the diskfile, everytime the file system is mounted. It returns -1 when disk open fails.

#### 2.1.3 dev\_close

This function closes the diskfile when the file system is unmounted.

#### 2.1.4 bio\_read

This function reads entries from the block and its inputs are const int block\_num and void \*buf. Based on the block\_num, this method identifies the block which is required, and calls pread function to read the data from the block in diskfile into buf. pread returns a status, which can be utilized to identify, if the block is empty or block read has failed. In case of empty, we return 0; if failed, we return -1.

#### 2.1.5 bio\_write

This function inputs are const int block\_num and void \*buf. Based on the block\_num, this method identifies the block which is required, and calls pwrite function to write the data from buf into the block in the diskfile. We return -1, if the program encounters a write error.

### 2.2 Inode

Inodes are crucial, as they function as the blueprint of the file system. All the data related file and directory(see next section) entries will be stored. We utilize the *struct inode* to handle the metadata of the entities. These would be really helpful in establishing any sort of relation between the directories, and upholds the tree-structure of the directories and files by maintaining a mapping between them.

### 2.2.1 readi

This function functions similar to the *bio\_read()*, it takes in the inode number *uint16\_t ino* and a inode pointer. It further uses the *s\_block* to identify the block in the disk, where the inode data is present. It finds the offset address within the block. Using the block, we call the *bio\_read()*. We set the inode pointer to the result we fetched from the block read.

### 2.2.2 writei

The inode counterpart of blocks, it takes in the inode number *uint16\_t ino* and a inode pointer and finds the block in the disk similar to *readi()*. Then we read the current data in the block, update the temporary node with the data from inode and do a block write, thereby updating the inode value.

## 2.3 Bitmap

Bitmaps are essential in both, inodes and blocks. The availability of Inodes and Blocks are resolved using bitmaps. There are five methods that rely primarily on these bitmap

### 2.3.1 get\_bitmap

Takes the bitmap(*ino/blk bitmap*) and the inode/block number as inputs, estimates the char bit to get based on the given number.

### 2.3.2 set\_bitmap

Takes the bitmap(*ino/blk bitmap*) and the inode/block number as inputs, estimates the char bit to set based on the given number.

### 2.3.3 unset\_bitmap

Takes the bitmap(*ino/blk bitmap*) and the inode/block number as inputs, estimates the char bit to unset based on the given number.

### 2.3.4 get\_avail\_ino

The function traverses the entire bitmap, from 0 to *max\_inum*. For every char byte, we use built in *ctz* to compute the trailing zero's position. If it is available, indicating there are free inode(s). In that case, we'll use the position which was returned by that function to update the bitmap inode and return this new position. If no inode is free, we return -1.

### 2.3.5 get\_avail\_blkno

The function is the block counterpart for *get\_avail\_ino()*, very similarly traverses the entire bitmap, from 0 to *max\_dnum*. For every char byte, we use built in *ctz* to compute the trailing zero's position. If it is available, indicating there are free block(s). In that case, we'll use the position which was returned by that function to update the bitmap block and return this new position. If no block is free, we return -1.

## 2.4 Files and Directories

### 2.4.1 dir\_find

This helper function will help us identify if the given folder/file name exists on any block/dirent of the given parent directory. The parent details, i.e., inode is retrieved using the inode number.

### 2.4.2 dir\_add

This method is similar to earlier, except that if this function finds the same file name, it exists. Otherwise, it will try to find an empty block and add the entry to our parent node.

### 2.4.3 `get_node_by_path`

Also called the `namei` operation, this takes the path name and iterates through the path till we reach the end, or throws an `ENOENT` error. We use string tokens on the path, and split the names of directories/file using a delimiter (`'/'`). Using the inode number, we use `dir_find` to iterate through all the entries of the given directory and comparing that against the current token, i.e., path. Finally if everything results in a success, we load the inode from memory using the final inode number.

### 2.4.4 `dir_remove`

The function is responsible for removing the inode entry from the parent inode. Subsequently, all the block related details will be updated, i.e, this will also free any unused memory blocks if the current block of memory is empty (`bio_write()`).

## 2.5 FUSE handlers

### 2.5.1 `rufs_init`

This function is invoked when the file system is mounted. We allocate memory for the superblocks and bitmaps which are necessary. If a diskfile already exists, we read the contents of this file and update all our primary data structures. Otherwise, this will call `rufs_mkfs()`.

### 2.5.2 `rufs_mkfs`

This function is the setup of our file system. We call the `dev_init()` to create a new diskfile. Later, we write all the information related to superblock on the disk. This is followed by initializing and writing the bitmap information onto the disk. Finally, we create the root directory and update the corresponding bitmap and inode.

### 2.5.3 `rufs_destroy`

This function stores the last known instance of all the data onto the disk, and de-allocates the data structures which are in memory. This includes the bitmaps, superblocks as well as calling `dev_close()` to close the diskfile.

### 2.5.4 `rufs_getattr`

This function takes the input path, finds the inode from the path using `get_node_by_path()`. If there's no inode found, we return the error code `ENOENT`(No such file or directory). Else, we update the `stbuf` with the stat of the directory/file inode.

### 2.5.5 `rufs_opendir`

This function returns 0 if it finds an entry in it's tree, otherwise it returns -1.

### 2.5.6 `rufs_readdir`

This function takes in path as an input, and finds the inode from the path using `get_node_by_path()`. Once found, it reads all the `dir_inode`'s data block and identifies the directory entry. It copies each directory entry into filler.

### 2.5.7 `rufs_mkdir`

This function takes in path, and uses `dirname()` and `basename()` to identify the parent and base names. First, it checks if parent exists, and return `ENOTDIR` if not found. Later, it attempts to add a new directory, by checking for free inodes. This is followed by a call to `dir_add()` to add directory entry of target directory to parent directory where it updates the parent directory inode. Finally, the inode is updated in the disk.

### 2.5.8 `rufs_create`

This function is similar to `rufs_mkdir()`, i.e, it attempts to create a new `dirent` entry, which is followed by adding a new file entry onto the disk.

### 2.5.9 `rufs_open`

This function takes the path, tries to find the inode entry from the path using `get_node_by_path()`, returns 0 if found, -1 if not.

### 2.5.10 `rufs_read`

This function takes the path, gets the inode entry from the path using `get_node_by_path()`. Based on the size and offset, it identifies the data block and reads it into the read buffer from the disk. We have ensured we only start reading the data from the offset into the buffer. Finally, we will return the total size that has been written onto the buffer.

### 2.5.11 `rufs_write`

This function takes the path, gets the inode entry from the path using `get_node_by_path()`. Based on the size and offset, it identifies the data block. We then write from the buffer to the disk. We have ensured we only start writing the data from the offset to disk. We then update the inode info of the base. Finally, we will return the total size that has been written into the disk.

### 2.5.12 `rufs_rmdir`

This function has been implemented similar to the linux implementation of `rmdir`. The folder will not be removed if there is any data inside the directory.

This function takes in path, and uses `dirname()` and `basename()` to identify the parent and base names. First, it checks if `dir_inode` exists for the parent and base, return `ENOTDIR` if not found. And then it unset the data block bitmap and inode bitmap in the disk. Also we reset the block and the inode to 0 in the disk. At the end, we call the `dir_remove()` to remove the directory entry of the base in its parent directory.

### 2.5.13 `rufs_unlink`

Very much similar to what `rufs_rmdir()` does, expect this method is called to delete files.

### 3 Benchmark testing

```

ks2020@kali:~/Documents/OS/project4$ ./final_out.sh
rm -f *.rufs
gcc -c *.c -o rufs.o -D_FILE_OFFSET_BITS=64 -Wno-discarded-qualifiers rufs.c -o rufs.o
gcc -c *.c -o rufs.o -D_FILE_OFFSET_BITS=64 -Wno-discarded-qualifiers block.c -o block.o
gcc -c *.c -o rufs.o -D_FILE_OFFSET_BITS=64 -Wno-discarded-qualifiers block.c -o block.o
rm -f simple_test test_cases
gcc -o simple_test test_cases.c
gcc -o test_cases test_cases.c
Simple test and reuse everything
TEST 1: File create Success
TEST 2: File write Success
TEST 3: File read Success
TEST 4: File read Success
TEST 5: Directory create Success
TEST 6: Sub-directory create Success
Benchmark completed with time 0.000689
Wednesday, May 01 2024: OPENING DISKFILE /common/home/ks2020/Documents/OS/project4/DISKFILE
Wednesday, May 01 2024: TOTAL BLOCKS INUSE AT OPEN 5
Wednesday, May 01 2024: TOTAL INODES INUSE AT OPEN 1
Wednesday, May 01 2024: TOTAL BLOCKS INUSE AT CLOSE 123
Wednesday, May 01 2024: BLOCKS USAGE HAS INCREASED BY 118
Wednesday, May 01 2024: TOTAL INODES INUSE AT CLOSE 183
Wednesday, May 01 2024: INODES USAGE HAS INCREASED BY 182
Wednesday, May 01 2024: CLOSING DISKFILE /common/home/ks2020/Documents/OS/project4/DISKFILE
total 12
drwxr-xr-x 2 ks2020 ks2020 4096 Apr 30 20:41 .
drwxr-xr-x 1 ks2020 ks2020 4096 Apr 30 20:41 ..
-rw-r--r-- 1 ks2020 ks2020 4096 May 1 04:29 file
drwxr-xr-x 11 ks2020 ks2020 13824 May 1 04:29 files
total 5
drwxr-xr-x 1 ks2020 ks2020 4096 May 1 04:29 .
drwxr-xr-x 1 ks2020 ks2020 4096 Apr 30 20:41 ..
Wednesday, May 01 2024: OPENING DISKFILE /common/home/ks2020/Documents/OS/project4/DISKFILE
Wednesday, May 01 2024: TOTAL BLOCKS INUSE AT OPEN 183
Wednesday, May 01 2024: TOTAL INODES INUSE AT OPEN 1
Wednesday, May 01 2024: BLOCKS USAGE HAS INCREASED BY 118
Wednesday, May 01 2024: TOTAL INODES INUSE AT CLOSE 1
Wednesday, May 01 2024: INODES USAGE HAS INCREASED BY 182
Wednesday, May 01 2024: CLOSING DISKFILE /common/home/ks2020/Documents/OS/project4/DISKFILE
Simple test and reuse everything
TEST 1: File create Success
TEST 2: File write Success
TEST 3: File read Success
TEST 4: File read Success
TEST 5: Directory create Success
TEST 6: Sub-directory create Success
Benchmark completed with time 0.000689
Wednesday, May 01 2024: OPENING DISKFILE /common/home/ks2020/Documents/OS/project4/DISKFILE
Wednesday, May 01 2024: TOTAL BLOCKS INUSE AT OPEN 5
Wednesday, May 01 2024: TOTAL INODES INUSE AT OPEN 1
Wednesday, May 01 2024: BLOCKS USAGE HAS INCREASED BY 570
Wednesday, May 01 2024: TOTAL INODES INUSE AT CLOSE 6
Wednesday, May 01 2024: INODES USAGE HAS INCREASED BY 514
Wednesday, May 01 2024: CLOSING DISKFILE /common/home/ks2020/Documents/OS/project4/DISKFILE
total 12
drwxr-xr-x 2 ks2020 ks2020 4096 May 1 04:29 .
drwxr-xr-x 1 ks2020 ks2020 4096 Apr 30 20:41 ..
-rw-r--r-- 1 ks2020 ks2020 4096 May 1 04:29 file
drwxr-xr-x 11 ks2020 ks2020 13824 May 1 04:29 files
total 5
drwxr-xr-x 1 ks2020 ks2020 4096 May 1 04:29 .
drwxr-xr-x 1 ks2020 ks2020 4096 Apr 30 20:41 ..
Wednesday, May 01 2024: OPENING DISKFILE /common/home/ks2020/Documents/OS/project4/DISKFILE
Wednesday, May 01 2024: TOTAL BLOCKS INUSE AT OPEN 575
Wednesday, May 01 2024: TOTAL INODES INUSE AT OPEN 515
Wednesday, May 01 2024: BLOCKS USAGE HAS INCREASED BY 569
Wednesday, May 01 2024: TOTAL INODES INUSE AT CLOSE 1
Wednesday, May 01 2024: INODES USAGE HAS INCREASED BY 514
Wednesday, May 01 2024: CLOSING DISKFILE /common/home/ks2020/Documents/OS/project4/DISKFILE
ks2020@kali:~/Documents/OS/project4$

```

Figure 1: simple\_test, test\_cases output

```

ks2020@kali:~/Documents/OS/project4$ ./rufs -s /tmp/ks2025/mountdir/ ; benchmark/simple_test ;
TEST 1: File create Success
TEST 2: File write Success
TEST 3: File close Success
TEST 4: File read Success
TEST 5: Directory create Success
TEST 6: Sub-directory create Success
Benchmark completed with time 0.000837
ks2020@kali:~/Documents/OS/project4$ findmnt | grep ks2025
| /tmp/ks2025/mountdir/ rufs fuse,rufs rw,nosuid,nodev,relative_user_id=1965237049,group_id=1965237049
ks2020@kali:~/Documents/OS/project4$ df -T | grep 1965237049
tmpfs 2839424 152 3289272 16 /run/user/1965237049
ks2020@kali:~/Documents/OS/project4$ fusermount -u /tmp/ks2025/mountdir; cat rufs_stats.log
Wednesday, May 01 2024: OPENING DISKFILE /common/home/ks2020/Documents/OS/project4/DISKFILE
Wednesday, May 01 2024: TOTAL BLOCKS INUSE AT OPEN 6
Wednesday, May 01 2024: TOTAL INODES INUSE AT OPEN 1
Wednesday, May 01 2024: TOTAL BLOCKS INUSE AT CLOSE 124
Wednesday, May 01 2024: BLOCKS USAGE HAS INCREASED BY 118
Wednesday, May 01 2024: TOTAL INODES INUSE AT CLOSE 183
Wednesday, May 01 2024: INODES USAGE HAS INCREASED BY 182
Wednesday, May 01 2024: CLOSING DISKFILE /common/home/ks2020/Documents/OS/project4/DISKFILE
ks2020@kali:~/Documents/OS/project4$

```

Figure 2: simple\_test case with block count

figure 1) tests both the cases. In our implementation, we accounted for "indirect\_ptrs" and increased the threshold of test\_cases. simple\_test took 0.000689 seconds to execute.

figure 2) shows the number of blocks that are being used by our disk file. We identified that the mount directory is using 152 blocks, when we run the command  $df -T <id>$ . But, when we print the same data onto our log file, we notice that it is only using 124 blocks instead. Please note, that these blocks also include the super block, and the number of blocks occupied by the inodes.

## 4 Challenges faced while implementing

### 4.1 Bitmaps

Instead of using bitmaps like we did earlier, we wanted to take a different route, and explored ways to reduce the total iterations made. We settled on `_builtin_ctz` for calculating the available bits for a given char, and `_builtin_popcount` to get the total bits used by that char.

### 4.2 Indirect pointers

The major roadblock I faced was identifying which block was freed and which wasn't. This was especially troublesome in `mkdir` & `rmdir`. Debugging was not easy, and the n number of print statements didn't help either. After going through the code n number of times, removing the direct pointers entirely, I could pin point the cause of the issue and fix it eventually. Similarly, identifying how to offset the read and write buffer was challenging since it was not as "direct" like earlier.